




**北京理工大学**  
**BEIJING INSTITUTE OF TECHNOLOGY**

**Computer Vision**



# Lecture 5 Edge Detection

---

**School of Computer Science and Technology**

**Ying Fu**



[fuying@bit.edu.cn](mailto:fuying@bit.edu.cn)



# Outline

---




- Overview
- Image gradients
- A simple edge detector
- Sobel edge detector
- Canny edge detector

# What are image edges?



grayscale image



Very sharp  
discontinuities  
in intensity.

domain  $\mathbf{x} = \begin{bmatrix} x \\ y \end{bmatrix}$

# Edges

---




# Edges



(A) Cave painting at Chauvet, France, about 30,000 B.C.;  
(B) Aerial photograph of the picture of a monkey as part of the Nazca Lines geoglyphs, Peru, about 700 – 200 B.C.;  
(C) Shen Zhou (1427-1509 A.D.): Poet on a mountain top, ink on paper, China;  
(D) Line drawing by 7-year old I. Lleras (2010 A.D.).

# Edges



Hubel & Wiesel, 1960s



Edge Detection



Overview





# Edges



- We know edges are special from human (mammalian) vision studies



# Edges





# Edge detection

---




- **Goal:** Identify sudden changes (discontinuities) in an image
  - Intuitively, most semantic and shape information from the image can be encoded in the edges
  - More compact than pixels
- **Ideal:** artist's line drawing (but artist is also using object-level knowledge)



# Edge detection



- Why do we care about edges?
  - Extract information, recognize objects
  - Recover geometry and viewpoint



# Origin of edges

---



- surface normal discontinuity
- depth discontinuity
- surface color discontinuity
- Illumination discontinuity

# Closeup of edges



Surface normal discontinuity



# Closeup of edges



Depth discontinuity





# Closeup of edges



Surface color discontinuity



# Closeup of edges



Illumination discontinuity



# Edge detection

---



- How would you go about detecting edges in an image (i.e., discontinuities in a function)?
  - ✓ You take derivatives: derivatives are large at discontinuities.
- How do you differentiate a discrete image (or any other discrete signal)?
  - ✓ You use (finite) differences.

# Outline

---



- Overview
- Image gradients
- A simple edge detector
- Sobel edge detector
- Canny edge detector

# Image gradients

---



- Derivatives in 1D

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x) = f_x$$



# Image gradients

---



- Derivatives in 1D — example

$$y = x^2 + x^4$$

$$\frac{dy}{dx} = 2x + 4x^3$$

$$y = \sin x + e^{-x}$$

$$\frac{dy}{dx} = \cos x + (-1)e^{-x}$$



# Image gradients

---

- Discrete Derivative in 1D

$$\frac{df}{dx} = \lim_{\Delta x \rightarrow 0} \frac{f(x) - f(x - \Delta x)}{\Delta x} = f'(x)$$

$$\frac{df}{dx} = \frac{f(x) - f(x - 1)}{1} = f'(x)$$

$$\frac{df}{dx} = f(x) - f(x - 1) = f'(x)$$



# Image gradients

---

- Types of Discrete derivative in 1D

Backward  $\frac{df}{dx} = f(x) - f(x-1) = f'(x)$

Forward  $\frac{df}{dx} = f(x) - f(x+1) = f'(x)$

Central  $\frac{df}{dx} = f(x+1) - f(x-1) = f'(x)$

# Image gradients

---



- 1D discrete derivate filters
  - Backward filter

$$f(x) - f(x-1) = f'(x) \quad [0 \ 1 \ -1]$$

- Forward filter

$$f(x) - f(x+1) = f'(x) \quad [-1 \ 1 \ 0]$$

- Central filter

$$f(x+1) - f(x-1) = f'(x) \quad [1 \ 0 \ -1]$$



# Image gradients

---

- Discrete derivate in 2D

- Given function  $f(x, y)$

- Gradient vector  $\nabla f(x, y) = \begin{bmatrix} \frac{\partial f(x, y)}{\partial x} \\ \frac{\partial f(x, y)}{\partial y} \end{bmatrix} = \begin{bmatrix} f_x \\ f_y \end{bmatrix}$

- Gradient magnitude  $|\nabla f(x, y)| = \sqrt{f_x^2 + f_y^2}$

- Gradient director  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$



# Image gradients



- 2D discrete derivative filters

What about this filter?

$$\frac{1}{3} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

- in what direction do  $x$  and  $y$  increase?



# Image gradients

- 2D discrete derivative — example

What happens when we apply this filter?

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



# Image gradients

- 2D discrete derivative — example

What happens when we apply this filter?

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



$$I_y = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$



# Image gradients

- 2D discrete derivative — example

Now let's try the other filter!

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$



# Image gradients

- 2D discrete derivative — example

What happens when we apply this filter?

$$I = \begin{bmatrix} 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \\ 10 & 10 & 20 & 20 & 20 \end{bmatrix}$$



$$I_x = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 10 & 10 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\frac{1}{3} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$



# Image gradients




- 3x3 image gradient filters

$$\frac{1}{3} \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$


$$\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



Derivative in x direction



Derivative in y direction



# Outline

---




- Overview
- Image gradients
- A simple edge detector
- Sobel edge detector
- Canny edge detector



# Characterizing edges



- An edge is a place of rapid change in the image intensity function






# Intensity profile





# Image gradient

- The gradient of an image:  $\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$



The gradient vector points in the direction of most rapid increase in intensity

The gradient direction is given by  $\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$

- how does this relate to the direction of the edge?

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

# Discrete derivative/gradient: example





- Which one is the gradient in the x-direction?  
How about y-direction?

Original  
Image



Gradient  
magnitude







# Effects of noise





- Consider a single row or column of the image
  - Plotting intensity as a function of position gives a signal



Where is the edge?

# Effects of noise

---



# Effects of noise

---



- Discrete gradient filters respond strongly to noise
  - Image noise results in pixels that look very different from their neighbors
  - Generally, the larger the noise the stronger the response
- What is to be done?
  - Smoothing the image should help, by forcing pixels different to their neighbors (=noise pixels?) to look more like neighbors

# Smoothing filters




- Mean smoothing

$$\begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 1 \end{bmatrix}$$

- Gaussian smoothing




$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

# A simple edge detector



- Smoothing + peaks



- To find edges, look for peaks in  $\frac{d}{dx}(f * g)$


# Derivative theorem of convolution



- This theorem gives us a very useful property


$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

- This saves us one operation






# Derivative of Gaussian (DOG) filter



2D-gaussian


\*

$$[1 \quad 0 \quad -1] =$$





x - derivative


# Derivative of Gaussian filter




x-direction



y-direction



# Tradeoff between smoothing and localization



- Stronger smoothing removes noise, but blurs edges.
- Finds edges at different “scales”.

# Laplace filter



Basically a second derivative filter.

- We can use finite differences to derive it, as with first derivative filter.

first-order  
finite difference

$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x + 0.5h) - f(x - 0.5h)}{h}$$

1D derivative filter

→ 

1	0	-1
---	---	----

second-order  
finite difference

$$f''(x) = \lim_{h \rightarrow 0} \frac{f(x + h) - 2f(x) + f(x - h)}{h^2}$$

Laplace filter

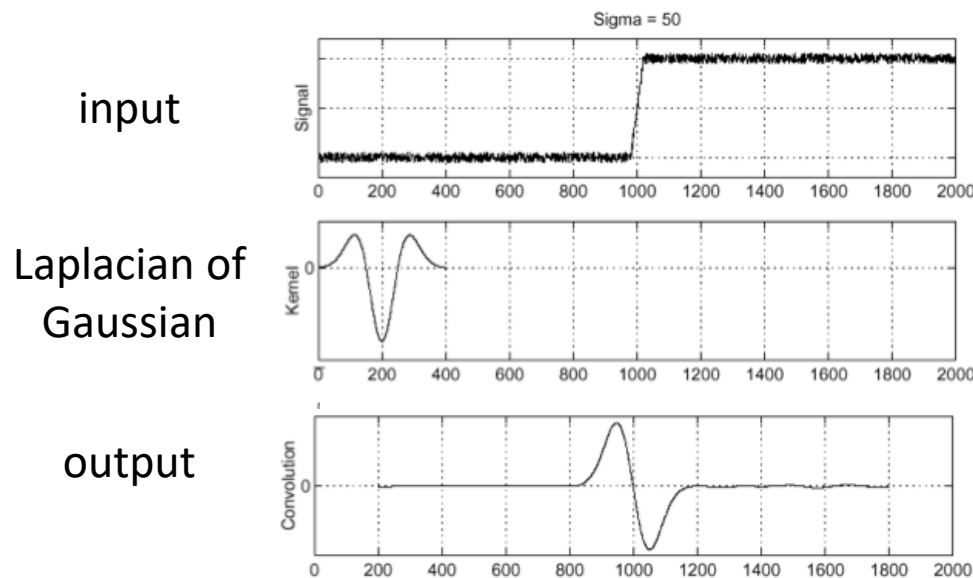
→ 

1	-2	1
---	----	---

# Laplacian of Gaussian (LoG) filter



- As with derivative, we can combine Laplace filtering with Gaussian filtering



“zero crossings” at edges



# Laplace and LoG filtering examples



Laplacian of Gaussian filtering



Laplace filtering



# Laplacian of Gaussian vs Derivative of Gaussian

---



Laplacian of Gaussian filtering




Derivative of Gaussian filtering


# Laplacian of Gaussian vs Derivative of Gaussian



- Zero crossings are more accurate at localizing edges (but not very convenient)




Laplacian of Gaussian filtering




Derivative of Gaussian filtering

# 2D Gaussian filters




$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$

Gaussian



$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$

Derivative of Gaussian




$$\nabla^2 h_{\sigma}(u, v)$$

Laplacian of Gaussian

# Designing an edge detector




- Criteria for an “optimal” edge detector:
  - (1) **Good detection:** the optimal detector must minimize the probability of false positives (detecting spurious edges caused by noise), as well as that of false negatives (missing real edges)



# Designing an edge detector



- Criteria for an “optimal” edge detector:  
**(2) Good localization:** the edges detected must be as close as possible to the true edges




# Designing an edge detector



- Criteria for an “optimal” edge detector:
  - (3) Single response:** the detector must return one point only for each true edge point; that is, minimize the number of local maxima around the true edge




True  
edge



Poor robustness  
to noise



Poor  
localization



Too many  
responses

# Outline

---



- Overview
- Image gradients
- A simple edge detector
- Sobel edge detector
- Canny edge detector
- Hough transform for line detection



# Sobel Operator



- Uses two  $3 \times 3$  kernels which are convolved with the original image to calculate approximations of the derivatives
- One for horizontal changes, and one for vertical

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} \quad \mathbf{G}_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

# Sobel Operation



- Smoothing + differentiation

$$\mathbf{G}_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \begin{bmatrix} +1 & 0 & -1 \end{bmatrix}$$

Gaussiansmoothing

differentiation

# The Sobel filter



- Horizontal Sober filter:

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$

- Vertical Sobel filter:

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} = \begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$



# Sobel Operation

---

- Magnitude:

$$\mathbf{G} = \sqrt{\mathbf{G}_x^2 + \mathbf{G}_y^2}$$


- Angle or direction of the gradient:

$$\Theta = \text{atan}\left(\frac{\mathbf{G}_y}{\mathbf{G}_x}\right)$$


# Sobel filter example



original



which Sobel filter?




which Sobel filter?

# Sobel filter example

---



original



horizontal Sobel filter




vertical Sobel filter



# Sobel filter example


---



original




horizontal Sobel filter



vertical Sobel filter



# Sobel Filter Problems



- Poor Localization (Trigger response in multiple adjacent pixels)
- Thresholding value favors certain directions over others
  - Can miss oblique edges more than horizontal or vertical edges
  - False negatives (missing real edges)

# Several derivative filters



Sobel

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

Scharr

3	0	-3
10	0	-10
3	0	-3

3	10	3
0	0	0
-3	-10	-3

Prewitt

1	0	-1
1	0	-1
1	0	-1

1	1	1
0	0	0
-1	-1	-1

Roberts

0	1
-1	0

1	0
0	-1

- How are the other filters derived and how do they relate to the Sobel filter?
- How would you derive a derivative filter that is larger than 3x3?



# Outline

---



- Overview
- Image gradients
- A simple edge detector
- Sobel edge detector
- Canny edge detector

# Canny edge detector

---



- This is probably the most widely used edge detector in computer vision
- Theoretical model: step-edges corrupted by additive Gaussian noise
- Canny has shown that the first derivative of the Gaussian closely approximates the operator that optimizes the product of signal-to-noise ratio and localization


# Canny edge detector

---




- Suppress Noise
- Compute gradient magnitude and direction
- Apply Non-Maximum Suppression
- Use hysteresis and connectivity analysis to detect edges

# Derivative of Gaussian (DOG) filter




x-direction




y-direction

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$



$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$



# Example

---




Original image




# Computing gradients

---



X-Derivative of Gaussian



Y-Derivative of Gaussian




Gradient Magnitude

# Get orientation at each pixel




$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$



# Compute gradients

---



Gradient  
Magnitude

# Canny edge detector

---



- Suppress Noise
- Compute gradient magnitude and direction
- Apply Non-Maximum Suppression

# Non-maximum suppression

---




- Edge occurs where gradient reaches a maxima
- Suppress non-maxima gradient even if it passes threshold
- Compare current pixel vs neighbors along direction of gradient
  - Remove if not maximum

# Remove spurious gradients




$|\nabla G|(x, y)$  is the gradient at pixel  $(x, y)$




$$M(x, y) = \begin{cases} |\nabla G|(x, y) & \text{if } |\nabla G|(x, y) > |\nabla G|(x', y') \\ & \& |\nabla G|(x, y) > |\nabla G|(x'', y'') \\ 0 & \text{otherwise} \end{cases}$$

$x'$  and  $x''$  are the neighbors of  $x$  along normal direction to an edge

# Non-maximum suppression



At  $q$ , we have a maximum if the value is larger than those at both  $p$  and at  $r$ . Interpolate to get these values.





# Non-maximum suppression

---



Before



After

# Canny edge detector

---



- Suppress Noise
- Compute gradient magnitude and direction
- Apply Non-Maximum Suppression
- Use hysteresis and connectivity analysis to detect edges

# Detecting edges with a single threshold

---



Threshold too high



Threshold too low




# Hysteresis thresholding

---

- Define two thresholds: Low and High
- If less than Low, not an edge
- If greater than High, strong edge
- If between Low and High, weak edge
  - Consider its neighbors iteratively then declare it an “edge pixel” if it is connected to an ‘strong edge pixel’ directly and via pixels between Low and High

# Hysteresis thresholding

---



# Final Canny Edges

---



# Canny edge detector

---



1. Filter image with **a filter**
2. Find magnitude and orientation of gradient
3. Non-maximum suppression:
  - Thin multi-pixel wide “ridges” down to single pixel width
4. Thresholding and linking (hysteresis):
  - Define two thresholds: low and high
  - Use the high threshold to start edge curves and the low threshold to continue them



# Effect of $\sigma$ (Gaussian kernel spread/size)



- The choice of  $\sigma$  depends on desired behavior
  - Large  $\sigma$  detects large scale edges
  - Small  $\sigma$  detects fine features



original



Canny with  $\sigma = 1$



Canny with  $\sigma = 2$

# References

---



- Basic reading:
  - Szeliski textbook, Chapter 3.2, 4,1